

Richard Mathematical Implementation

Rob Jinman

January 14, 2024

Dense Layers

Forward pass

The weighted sums z_i^l for layer l are given by

$$\boxed{z_i^l = \sum_j w_{i,j}^l a_j^{l-1} + b_i^l} \quad (1)$$

and we obtain the activations by applying the sigmoid function

$$\boxed{a_i^l = \sigma(z_i^l)} \quad (2)$$

where

$$\boxed{\sigma(x) = \frac{1}{1 + e^{-x}}} \quad (3)$$

In vectorized form, the forward pass can be written as

$$\boxed{\begin{aligned} \mathbf{z}^l &= \mathbf{W}^l \mathbf{a}^{l-1} + \mathbf{b}^l \\ \mathbf{a}^l &= \sigma(\mathbf{z}^l) \end{aligned}} \quad (4)$$

$$(5)$$

Backward pass

Given $\partial C / \partial a_i^l$ back-propagated from layer $l + 1$, we calculate the layer delta

$$\boxed{\delta_i^l := \frac{\partial C}{\partial z_i^l}} \quad (6)$$

Applying the chain rule gives

$$\frac{\partial C}{\partial z_i^l} = \frac{\partial C}{\partial a_i^l} \frac{\partial a_i^l}{\partial z_i^l}$$

where $\partial a_i^l / \partial z_i^l$ is the derivative of the sigmoid function, so

$$\boxed{\delta_i^l = \frac{\partial C}{\partial a_i^l} \sigma'(z_i^l)} \quad (7)$$

which in vectorized form is the hadamard product

$$\boxed{\boldsymbol{\delta}^l = \nabla_{\mathbf{a}^l} C \odot \sigma'(\mathbf{z}^l)} \quad (8)$$

We use this value to compute the cost gradient both with respect to the layer parameters and the layer inputs.

Cost gradient with respect to parameters

For the weights $w_{i,j}^l$, applying the chain-rule gives

$$\frac{\partial C}{\partial w_{i,j}^l} = \frac{\partial C}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{i,j}^l} = \delta_i^l \frac{\partial z_i^l}{\partial w_{i,j}^l} = \delta_i^l \frac{\partial \sum_k w_{i,k}^l a_k^{l-1} + b_i^l}{\partial w_{i,j}^l} = \delta_i^l a_j^{l-1}$$

which in vectorized form is the outer product between $\boldsymbol{\delta}^l$ and \mathbf{a}^{l-1}

$$\boxed{\nabla_{W^l} C = \boldsymbol{\delta}^l \otimes \mathbf{a}^{l-1}} \quad (9)$$

For the bias b_i^l we again use the chain rule

$$\frac{\partial C}{\partial b_i^l} = \frac{\partial C}{\partial z_i^l} \frac{\partial z_i^l}{\partial b_i^l} = \delta_i^l \frac{\partial z_i^l}{\partial b_i^l} = \delta_i^l \frac{\partial \sum_k w_{i,k}^l a_k^{l-1} + b_i^l}{\partial b_i^l} = \delta_i^l$$

which in vectorized form is

$$\boxed{\nabla_{b^l} C = \boldsymbol{\delta}^l} \quad (10)$$

Then, given a learning rate λ , we update the weights and biases

$$\boxed{\mathbf{W}^l \leftarrow \mathbf{W}^l - \lambda \nabla_{W^l} C} \quad (11)$$

$$\mathbf{b}^l \leftarrow \mathbf{b}^l - \lambda \nabla_{b^l} C \quad (12)$$

Cost gradient with respect to layer inputs

The multi-variable chain rule gives us

$$\frac{\partial C}{\partial a_i^{l-1}} = \sum_j \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial a_i^{l-1}} = \sum_j \delta_j^l \frac{\partial z_j^l}{\partial a_i^{l-1}} = \sum_j \delta_j^l \frac{\partial \sum_k w_{j,k}^l a_k^{l-1} + b_j^l}{\partial a_i^{l-1}} = \sum_j \delta_j^l w_{j,i}^l$$

which is the layer delta multiplied by the transposed weight matrix

$$\boxed{\nabla_{a^{l-1}} C = (\mathbf{W}^l)^T \boldsymbol{\delta}^l} \quad (13)$$

This value is propagated back to layer $l - 1$.

Output Layer

Forward pass

The forward pass is the same as for ordinary dense layers (see above).

Backward pass

It's at this final layer L where backpropagation begins. As in equation (8), we have

$$\boxed{\boldsymbol{\delta}^L = \nabla_{a^L} C \odot \sigma'(\mathbf{z}^L)} \quad (14)$$

But rather than receive $\nabla_{a^L} C$ from the next layer (there is no next layer), we compute it directly. The cost function $C(\mathbf{a}^L, \mathbf{y})$ computes the squared error between the network output \mathbf{a}^L and the expected network output \mathbf{y} for the current sample.

$$\boxed{C(\mathbf{a}^L, \mathbf{y}) = \frac{1}{2} \|\mathbf{y} - \mathbf{a}^L\|^2} \quad (15)$$

which in component form is

$$\frac{1}{2} \sum_i (a_i^L - y_i)^2 = \frac{1}{2} \sum_i \left((a_i^L)^2 - 2a_i^L y_i + (y_i)^2 \right)$$

Differentiating with respect to a_i^L we get

$$\boxed{\frac{\partial C}{\partial a_i^L} = a_i^L - y_i} \quad (16)$$

which in vectorized form is

$$\boxed{\nabla_{\mathbf{a}^L} C = \mathbf{a}^L - \mathbf{y}} \quad (17)$$

which we plug into equation (14) to obtain δ^L . Using this value we calculate the gradients for the layer parameters and inputs in the same way as for the dense layers.

Convolutional Layers

The convolution between functions f and g is

$$\boxed{(f * g)(t) := \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau} \quad (18)$$

The discrete form in 2 dimensions is

$$\boxed{(f * g)(x, y) := \sum_i \sum_j f(i, j)g(x - i, y - j)} \quad (19)$$

which is equivalent to cross-correlation with a horizontally and vertically flipped kernel. The cross-correlation is defined as

$$\boxed{(f \star g)(x, y) := \sum_i \sum_j f(i, j)g(x + i, y + j)} \quad (20)$$

Forward pass

A layer of depth D generates D feature maps using D kernels and biases. During the forward pass we compute the elements of the feature map $z_{x,y}^{l,d}$ by cross-correlating the 3-dimensional block of input activations \mathbf{A}^{l-1} with the 3-dimensional kernel $\mathbf{W}^{l,d}$ and adding the bias $b^{l,d}$. We then obtain the activations by applying the ReLU function.

$$\boxed{\mathbf{Z}^{l,d} = (\mathbf{A}^{l-1} \star \mathbf{W}^{l,d}) + b^{l,d}} \quad (21)$$

$$\mathbf{A}^{l,d} = R(\mathbf{Z}^{l,d}) \quad (22)$$

where ReLU is defined as

$$\boxed{R(z) = \max(0, z)} \quad (23)$$

In component form, the forward pass looks like

$$\boxed{z_{x,y}^{l,d} = \sum_i \sum_j \sum_k w_{i,j,k}^{l,d} a_{x+i,y+j,k}^{l-1} + b^{l,d}} \quad (24)$$

$$a_{x,y}^{l,d} = R(z_{x,y}^{l,d}) \quad (25)$$

Backward pass

As before, we calculate the layer delta

$$\delta_{x,y}^{l,d} := \frac{\partial C}{\partial z_{x,y}^{l,d}}$$

Applying the chain rule, we get

$$\delta_{x,y}^{l,d} = \frac{\partial C}{\partial a_{x,y}^{l,d}} R'(z_{x,y}^{l,d})$$

where $\partial C / \partial a_{x,y}^{l,d}$ is propagated back from layer $l + 1$, and R' is the ReLU derivative.

$$\boxed{R'(z) = \begin{cases} 0 & \text{if } z \leq 0 \\ 1 & \text{if } z > 0 \end{cases}} \quad (26)$$

Cost gradient with respect to layer inputs

From the chain rule, we have

$$\frac{\partial C}{\partial a_{x,y,z}^{l-1}} = \sum_d \left(\sum_{x'} \sum_{y'} \frac{\partial C}{\partial z_{x',y'}^{l,d}} \frac{\partial z_{x',y'}^{l,d}}{\partial a_{x,y,z}^{l-1}} \right)$$

which is

$$\frac{\partial C}{\partial a_{x,y,z}^{l-1}} = \sum_d \left(\sum_{x'} \sum_{y'} \delta_{x',y'}^{l,d} \frac{\partial z_{x',y'}^{l,d}}{\partial a_{x,y,z}^{l-1}} \right)$$

Expanding the $z_{x',y'}^{l,d}$ term

$$\frac{\partial C}{\partial a_{x,y,z}^{l-1}} = \sum_d \left(\sum_{x'} \sum_{y'} \delta_{x',y'}^{l,d} \frac{\partial \sum_i \sum_j \sum_k w_{i,j,k}^{l,d} a_{x'+i,y'+j,k}^{l-1} + b^{l,d}}{\partial a_{x,y,z}^{l-1}} \right)$$

All terms of the differential vanish except where $x = x' + i$, $y = y' + j$, and $z = k$, or equivalently $i = x - x'$, $j = y - y'$, and $k = z$, therefore

$$\frac{\partial C}{\partial a_{x,y,z}^{l-1}} = \sum_d \left(\sum_{x'} \sum_{y'} \delta_{x',y'}^{l,d} w_{x-x',y-y',z}^{l,d} \right)$$

which amounts to a *full* 2-dimensional convolution between each 2-dimensional kernel slice $\mathbf{W}_z^{l,d}$ and feature map delta $\boldsymbol{\delta}^{l,d}$, all added together.

$$\boxed{\nabla_{\mathbf{a}_z^{l-1}} C = \sum_d (\mathbf{W}_z^{l,d} * \boldsymbol{\delta}^{l,d})} \quad (27)$$

Cost gradient with respect to layer parameters

From the chain rule, we have

$$\frac{\partial C}{\partial w_{i,j,k}^{l,d}} = \sum_x \sum_y \frac{\partial C}{\partial z_{x,y}^{l,d}} \frac{\partial z_{x,y}^{l,d}}{\partial w_{i,j,k}^{l,d}} = \sum_x \sum_y \delta_{x,y}^{l,d} \frac{\partial z_{x,y}^{l,d}}{\partial w_{i,j,k}^{l,d}}$$

Expanding $z_{x,y}^{l,d}$, we get

$$\frac{\partial C}{\partial w_{i,j,k}^{l,d}} = \sum_x \sum_y \delta_{x,y}^{l,d} \frac{\partial \sum_{i'} \sum_{j'} \sum_{k'} w_{i',j',k'}^{l,d} a_{x+i',y+j',k'}^{l-1} + b^{l,d}}{\partial w_{i,j,k}^{l,d}}$$

All terms in the differential vanish except where $i = i'$, $j = j'$, and $k = k'$, so finally

$$\frac{\partial C}{\partial w_{i,j,k}^{l,d}} = \sum_x \sum_y \delta_{x,y}^{l,d} a_{x+i,y+j,k}^{l-1}$$

which is the cross-correlation of 2-dimensional input slice \mathbf{A}_k^{l-1} with 2-dimensional feature map delta $\delta^{l,d}$

$$\boxed{\nabla_{\mathbf{w}_k^{l,d}} C = \mathbf{A}_k^{l-1} \star \delta^{l,d}} \quad (28)$$

Max Pooling Layers

Max pooling layers perform a downscale by allowing through only the largest values within each $M \times N$ window.

Forward pass

For $0 \leq i < M$ and $0 \leq j < N$, the elements of the max pooling layer are given by

$$\boxed{z_{x,y,z}^l = \max_{i,j} (a_{x+i,y+j,z}^{l-1})} \quad (29)$$

$$a_{x,y,z}^l = z_{x,y,z}^l \quad (30)$$

and we don't apply an activation function—or you could say the activation function is just the identity function $f(z) = z$.

Backward pass

As usual, the layer delta is

$$\delta_{x,y,z}^l := \frac{\partial C}{\partial z_{x,y,z}^l} = \frac{\partial C}{\partial a_{x,y,z}^l} \frac{\partial a_{x,y,z}^l}{\partial z_{x,y,z}^l}$$

There's no activation function, so $z_{x,y,z}^l$ and $a_{x,y,z}^l$ are identical, therefore

$$\boxed{\delta_{x,y,z}^l = \frac{\partial C}{\partial a_{x,y,z}^l}} \quad (31)$$

Or in vectorized form

$$\boxed{\delta^l = \nabla_{\mathbf{A}^l} C} \quad (32)$$

which is the value we receive from layer $l + 1$.

Cost gradient with respect to layer inputs

Each element $a_{x,y,z}^{l-1}$ only contributes to a single element $z_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l$ so there's no need for a summation.

$$\begin{aligned} \frac{\partial C}{\partial a_{x,y,z}^{l-1}} &= \frac{\partial C}{\partial z_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l} \frac{\partial z_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l}{\partial a_{x,y,z}^{l-1}} = \delta_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l \frac{\partial z_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l}{\partial a_{x,y,z}^{l-1}} \\ \frac{\partial C}{\partial a_{x,y,z}^{l-1}} &= \delta_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l \frac{\partial \max_{i,j}(a_{\lfloor x/M \rfloor + i, \lfloor y/N \rfloor + j, z}^{l-1})}{\partial a_{x,y,z}^{l-1}} \\ \boxed{\frac{\partial C}{\partial a_{x,y,z}^{l-1}} &= \begin{cases} \delta_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l & \text{if } a_{x,y,z}^{l-1} = \max_{i,j}(a_{\lfloor x/M \rfloor + i, \lfloor y/N \rfloor + j, z}^{l-1}) \\ 0 & \text{otherwise} \end{cases}} \end{aligned} \quad (33)$$

So each element $\partial C / \partial a_{x,y,z}^{l-1}$ is equal to $\delta_{\lfloor x/M \rfloor, \lfloor y/N \rfloor, z}^l$, but only if its corresponding value in \mathbf{A}^{l-1} was the largest within its window during the forward pass.